

Today you'll learn how to:

- [Load datasets](#)
- [Scrape Webpages](#)
- [Build REST APIs](#)
- [Analyze Data and Show Statistical Summaries](#)
- [Visualize Data](#)
- [Train a Machine Learning Model](#)
- [Develop Simple Web Applications](#)

Load datasets

To perform any sort of analysis, you first have to load in the data. With R, you can connect to any data so

For a simple demonstration, we'll see how to load in CSV data. You can find the Iris dataset in CSV form:

```
iris <- read.csv("iris.csv")
head(iris)
```

And here's what the `head` function outputs – the first six rows:

```
> head(iris)
  sepal.length sepal.width petal.length petal.width variety
1          5.1         3.5         1.4         0.2   Setosa
2          4.9         3.0         1.4         0.2   Setosa
3          4.7         3.2         1.3         0.2   Setosa
4          4.6         3.1         1.5         0.2   Setosa
5          5.0         3.6         1.4         0.2   Setosa
6          5.4         3.9         1.7         0.4   Setosa
```

Image 1 – Iris dataset head

Did you know there's no need to download the dataset? You can load it from the web:

```
iris <- read.csv("https://gist.githubusercontent.com/netj/8836201/raw/6f9306ad21398ea43cba4f7d0e07d5ae3/iris.csv")
head(iris)
```

That's all great, but what if you can't find an appropriate dataset? That's where web scraping comes into

Web scraping

A good dataset is difficult to find, so sometimes you have to be creative. Web scraping is considered as c

In R, the `rvest` package is used for the task. As some websites have strict policies against scraping, we

```
library(rvest)
```

```
url <- "http://books.toscrape.com/catalogue/category/books/travel_2/index.html"
titles <- read_html(url) %>%
  html_nodes("h3") %>%
  html_nodes("a") %>%
  html_text()
```

The `titles` variable contains the following elements:

```
> titles
[1] "It's Only the Himalayas"      "Full Moon over Noah's ..." "See America: A Celebration ..."
[4] "Vagabonding: An Uncommon Guide ..." "Under the Tuscan Sun"        "A Summer In Europe"
[7] "The Great Railway Bazaar"     "A Year in Provence ..."    "The Road to Little ..."
[10] "Neither Here nor There: ..." "1,000 Places to See ..."    "
```

Image 2 – Web Scraping example in R

Yes – it's that easy. Just don't cross any boundaries. Check if a website has a public API first – if so, there

Build REST APIs

With practical machine learning comes the issue of model deployment. Currently, the best option is to write

In R, the `plumber` package is used to build REST APIs. Here's the one that comes in by default when you

```
library(plumber)
```

```
## @apiTitle Plumber Example API
```

```
## Echo back the input
```

```
## @param msg The message to echo
```

```
## @get /echo
```

```
function(msg = "") {
  list(msg = paste0("The message is: '", msg, "'"))
}
```

```
## Plot a histogram
```

```
## @png
```

```
## @get /plot
```

```
function() {
  rand <- rnorm(100)
  hist(rand)
}
```

```
## Return the sum of two numbers
```

```
## @param a The first number to add
```

```
## @param b The second number to add
```

```
## @post /sum
```

```
function(a, b) {
  as.numeric(a) + as.numeric(b)
}
```

The API has three endpoints:

1. `/echo` – returns a specified message in the response
2. `/plot` – shows a histogram of 100 random normally distributed numbers
3. `/sum` – sums two numbers

The `plumber` package comes with Swagger UI, so you can explore and test your API in the web browser

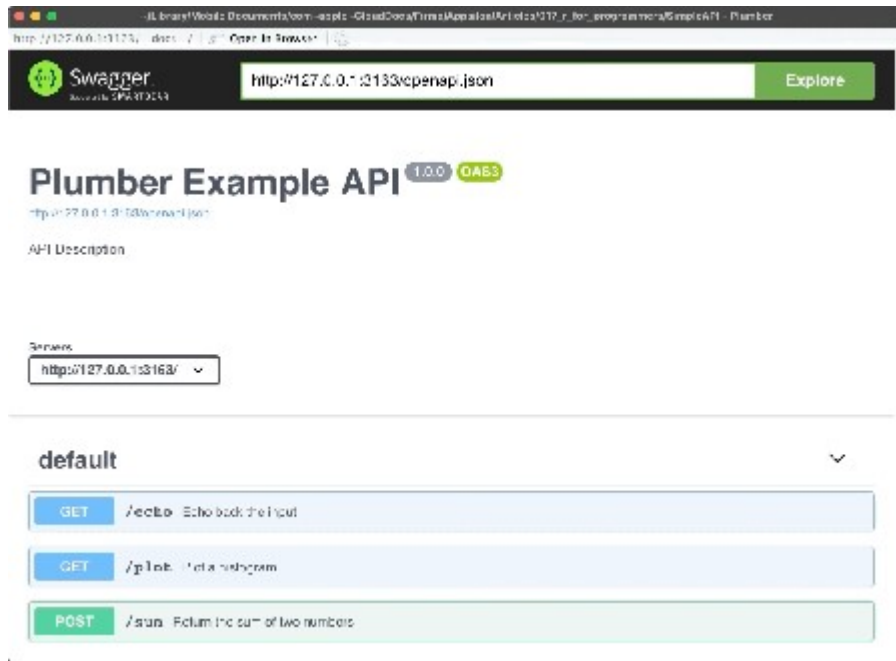


Image 3 – Plumber REST API Showcase

Statistics and Data Analysis

This is one of the biggest reasons why R is so popular. There are entire books and courses on this topic,

Most of the data manipulation in R is done with the `dplyr` package. Still, we need a dataset to manipulate

```
library(dplyr)
library(gapminder)
```

```
head(gapminder)
```

You should see the following in the console:

```
> head(gapminder)
# A tibble: 6 x 6
  country    continent  year lifeExp    pop gdpPercap
  <fct>      <fct>    <int>  <dbl>    <int>    <dbl>
1 Afghanistan Asia      1952   28.8  8425333    779.
2 Afghanistan Asia      1957   30.3  9240934    821.
3 Afghanistan Asia      1962   32.0 10267083    853.
4 Afghanistan Asia      1967   34.0 11537966    836.
5 Afghanistan Asia      1972   36.1 13079460    740.
6 Afghanistan Asia      1977   38.4 14880372    786.
```

Image 4 – Head of Gapminder dataset

To perform any kind of statistical analysis, you could use R's built-in functions such as `min`, `max`, `range`,

```
summary(gapminder)
```

Here's a statistical summary of the Gapminder dataset:

country	continent	year	lifeExp	pop	gdpPercap
Afghanistan: 12	Africa :624	Min. :1952	Min. :23.60	Min. :6.001e+04	Min. : 241.2
Albania : 12	Americas:300	1st Qu.:1966	1st Qu.:48.20	1st Qu.:2.794e+06	1st Qu.: 1202.1
Algeria : 12	Asia :396	Median :1980	Median :60.71	Median :7.024e+06	Median : 3531.8
Angola : 12	Europe :360	Mean :1980	Mean :59.47	Mean :2.960e+07	Mean : 7215.3
Argentina : 12	Oceania : 24	3rd Qu.:1993	3rd Qu.:70.85	3rd Qu.:1.959e+07	3rd Qu.: 9325.5
Australia : 12		Max. :2007	Max. :82.60	Max. :1.319e+09	Max. :113523.1
(Other) :1632					

Image 5 – Statistical summary of the Gapminder dataset

With `dplyr`, you can drill down and keep only the data of interest. Let's see how to show only data for Poland

```
gapminder %>%
  filter(continent == "Europe", country == "Poland") %>%
  mutate(TotalGDP = pop * gdpPercap)
```

The corresponding results are shown in the console:

	country	continent	year	lifeExp	pop	gdpPercap	TotalGDP
	<fct>	<fct>	<int>	<dbl>	<int>	<dbl>	<dbl>
1	Poland	Europe	1952	61.3	25730551	4029.	103676873316.
2	Poland	Europe	1957	65.8	28235346	4734.	133673272043.
3	Poland	Europe	1962	67.6	30329617	5339.	161922307755.
4	Poland	Europe	1967	69.6	31785378	6557.	208421579589.
5	Poland	Europe	1972	70.8	33039545	8007.	264531348088.
6	Poland	Europe	1977	70.7	34621254	9508.	329183780347.
7	Poland	Europe	1982	71.3	36227381	8452.	306176833715.
8	Poland	Europe	1987	71.0	37740710	9082.	342774381701.
9	Poland	Europe	1992	71.0	38370697	7739.	296946267448.
10	Poland	Europe	1997	72.8	38654957	10160.	392718270288.
11	Poland	Europe	2002	74.7	38625976	12002.	463598198650.
12	Poland	Europe	2007	75.6	38518241	15390.	592792827796.

Image 6 – History data and total GDP for Poland

Data Visualization

R is known for its impeccable data visualization capabilities. The `ggplot2` package is a good starting point

To start, we will create a line chart comparing the total population in Poland over time. We will need to filter

```
library(dplyr)
library(gapminder)
library(scales)
library(ggplot2)

poland <- gapminder %>%
  filter(continent == "Europe", country == "Poland")

ggplot(poland, aes(x = year, y = pop)) +
  geom_line(size = 2, color = "#0099f9") +
  ggtitle("Poland population over time") +
  xlab("Year") +
  ylab("Population") +
  expand_limits(y = c(10^6 * 25, NA)) +
  scale_y_continuous(
    labels = paste0(c(25, 30, 35, 40), "M"),
    breaks = 10^6 * c(25, 30, 35, 40)
  ) +
```

```
theme_bw()
```

Here is the corresponding output:




Image 7 – Poland population over time

You can get a similar visualization with the first two code lines – the others are added for styling.

The `ggplot2` package can display almost any data visualization type, so let's explore bar charts next. W

```
europe_2007 <- gapminder %>%  
  filter(continent == "Europe", year == 2007)  
  
ggplot(europe_2007, aes(x = reorder(country, -lifeExp), y = lifeExp)) +  
  geom_bar(stat = "identity", fill = "#0099f9") +  
  geom_text(aes(label = lifeExp), color = "white", hjust = 1.3) +  
  ggtitle("Average life expectancy in Europe countries in 2007") +  
  xlab("Country") +  
  ylab("Life expectancy (years)") +  
  coord_flip() +  
  theme_bw()
```

Here's how the chart looks like:




Image 8 – Average life expectancy in European countries in 2007

Once again, the first two code lines for the visualization will produce similar output. The rest are here to n

Training a Machine Learning Model

Yet another area that R handles with ease. The `rpart` package is great for machine learning, and we will

Here's how to load in the libraries, perform the train/test split, fit and visualize the model:

```
library(caTools)
library(rpart)
library(rpart.plot)

set.seed(42)
sample <- sample.split(iris, SplitRatio = 0.75)
iris_train = subset(iris, sample == TRUE)
iris_test = subset(iris, sample == FALSE)

model <- rpart(Species ~., data = iris_train, method = "class")
rpart.plot(model)
```

The snippet shouldn't take more than a second or two to execute. Once done, you'll be presented with th




Image 9 – Decision tree visualization for Iris dataset

The above figure tells you everything about the decision-making process of the algorithm. We can now e

```

preds <- predict(model, iris_test, type = "class")

confusion_matrix <- table(iris_test$Species, preds)
print(confusion_matrix)

accuracy <- sum(diag(confusion_matrix)) / sum(confusion_matrix)
print(accuracy)

> print(confusion_matrix)
      preds
      setosa versicolor virginica
setosa      20         0         0
versicolor  0         18         2
virginica   0          1        19
> print(accuracy)
[1] 0.95

```

Image 10 – Confusion matrix and accuracy on the test subset

As you can see, we got a 95% accurate model with only a couple of lines of code.

Develop Simple Web Applications

At [Appsilon](#), we are global leaders in R Shiny, and we've developed some of the world's most [advanced](#) f

For the web app example, we'll see how to make simple interactive dashboards that displays a scatter pl

Here is a script for the Shiny app:

```

library(shiny)
library(ggplot2)

ui <- fluidPage(
  sidebarPanel(
    width = 3,
    tags$h4("Select"),
    varSelectInput(
      inputId = "x_select",
      label = "X-Axis",
      data = mtcars
    ),
    varSelectInput(
      inputId = "y_select",
      label = "Y-Axis",
      data = mtcars
    )
  ),
  mainPanel(
    plotOutput(outputId = "scatter")
  )
)

server <- function(input, output) {

```

```

output$scatter <- renderPlot({
  col1 <- sym(input$x_select)
  col2 <- sym(input$y_select)

  ggplot(mtcars, aes(x = !!col1, y = !!col2)) +
    geom_point(size = 6, color = "#0099f9") +
    ggtitle("MTCars Dataset Explorer") +
    theme_bw()
})
}

shinyApp(ui = ui, server = server)

```

And here's the corresponding Shiny app:




Image 11 – MTCars Shiny app

This dashboard is as simple as they come, but that doesn't mean you can't develop beautiful-looking app

Looking for inspiration? [Take a look at our Shiny App Demo Gallery.](#)

Conclusion

To conclude – R can do almost anything that a general-purpose programming language can do. The que: